



# Module 9

*The Node*



The Corda logo, consisting of the word "corda" in a white sans-serif font, with a small red dot preceding the letter "c". The entire logo is contained within a red rectangular box.




# Learning outcomes

- Understand how the node is structured
- Learn what services are available to the node internally
- Learn what APIs are available to interact with the node
- Learn how to deploy and launch nodes




r3.

# What is a Node?



# Corda networks

- A Corda network is made up of nodes:




r3.

# Defining a node

- A node is an instance of the Corda software with a unique identity on the network
- Each node's identity is provided by a certificate signed by a network root authority
- Networks can decide what constitutes a valid identity – legal names, usernames, IP addresses...




# Node internals




r3.

# ServiceHub vs. CordaRPCOps

- **CordaRPCOps** defines how the owner interacts with their node
- **ServiceHub** defines how the node accesses its services internally





# The vault

- The node stores data in two ways:
  - **The vault**, which represents the node's personal ledger:
    - **Consumed states**, which represent an immutable ledger of the owner's past activity
    - **Unconsumed states**, which represent the owner's "balance"
    - **Attachments**, binary files linked to a transaction
  - **Local storage** for all other data (eg. transactions, metadata) that is needed for the node to interact with the ledger





# Default Node Services

# The ServiceHub interface

- **ServiceHub** represents the operations available internally to the node:

```
interface ServiceHub {  
    val myInfo: NodeInfo  
    val identityService: IdentityService  
    val attachments: AttachmentStorage  
    val vaultService: VaultService  
    val keyManagementService: KeyManagementService  
    val validatedTransactions: TransactionStorage  
    val networkMapCache: NetworkMapCache  
    val clock: Clock
```



# The ServiceHub

- The key operations provided by the **ServiceHub** can be divided into four types:
  1. Information on the node's identity and keys
  2. Information on other network nodes
  3. Retrieving the current time and scheduling events
  4. Retrieving or recording data in the vault or local storage





# The Observable pattern

- Many **ServiceHub** and **CordaRPCOps** methods return an instance of the **rx.Observable** class
- The observable pattern is as follows:
  - The observable emits events
  - An observer can subscribe to the observable
  - The observer will then be notified of any emitted events
- The observable pattern allows the node's owner to be automatically notified of updates to things such as the network map, the vault, and in-progress flows



# 1. Node information – myInfo

- Information on the node is obtained from **ServiceHub.myInfo**, which returns a **NodeInfo** instance:

```
data class NodeInfo(  
    val addresses: List<NetworkHostAndPort>,  
    val legalIdentities: List<Party>,  
    val legalIdentitiesAndCerts: Set<PartyAndCertificate>,  
    val platformVersion: Int,  
    val serial: Long)
```

- Where:
  - addresses** are the Artemis MQ address that allows other nodes to send the node messages
  - legalIdentitiesAndCerts** allows other nodes to identify this node

## 2. Network information – `networkMapCache`

- Communication between network nodes is point-to-point (i.e. between named nodes)
- To look up other nodes, the node uses `ServiceHub.networkMapCache`, which implements `NetworkMapCache`
- `NetworkMapCache` tracks four kinds of nodes:
  - `partyNodes`: every node on the network
  - `networkMapNodes`: nodes advertising a network map service
  - `notaryNodes`: nodes advertising a notary service
  - `regulatorNodes`: nodes advertising a regulatory service

r3.

## 2. Network information – networkMapCache

- The **networkMapCache** provides methods for retrieving nodes based on:
  - Name
  - Public key
  - Service advertised
- The **track()** method allows the node to learn about changes to the network map over time:

```
fun track(): DataFeed<List<NodeInfo>, MapChange>
```



## 2. Network information – Looking Up Nodes

- Some `networkMapCache` methods to look up nodes:

```
fun getNodeByLegalName(principal: CordaX500Name): NodeInfo?
```

```
fun getNodeByLegalIdentityKey(identityKey: PublicKey): NodeInfo?
```

```
fun getNodeByLegalIdentity(party: AbstractParty): NodeInfo?
```

```
fun getNodeByAddress(address: NetworkHostAndPort): NodeInfo?
```





### 3. Time information – clock

- In Corda, there is no network time. Each node tracks time separately
- Each node tracks time using `ServiceHub.clock`, which embeds an instance of Java's built-in `Java.Clock`
- This clock is used to choose time-stamps for transactions



## 4. Data-related services – The Vault

- The vault is accessed via `ServiceHub.vaultService`
- A rich API is available for querying the vault:  
<https://docs.corda.net/api-vault-query.html>
- A simple example: here is how we could extract all `IOUState`s:

```
serviceHub  
    .vaultService  
    .queryBy<IOUState>(queryCriteria)  
    .states
```



## 4. Data-related services – Storage

- The node's local storage is only accessed using:
  - **validatedTransactions**, which provides methods to access the node's stored transactions



r3.

# Interacting with the Node

# Interacting with the node

- A node's owner interacts with the node using the **CordaRPCOps** interface
- The key operations in **CordaRPCOps** are:
  - **startTrackedFlowDynamic** to start flows
  - **vaultQueryBy** to access the vault
  - Methods to check for, upload and open attachments



# Starting flows


- You command a node to start a flow using:

```
fun <T : Any> startTrackedFlowDynamic(  
    logicType: Class<out FlowLogic<T>>, vararg args: Any?)  
    : FlowHandle<T>
```

- This method returns a **FlowProgressHandle**:

```
data class FlowProgressHandle<A>(  
    val id: StateMachineRunId,  
    val progress: Observable<String>,  
    val returnValue: ListenableFuture<A>)
```

- Where:
  - **progress** is an **Observable** allowing you to track the flow's progress
  - **returnValue** is a **ListenableFuture** representing the asynchronous result of running the flow



# Accessing the vault

- The node owner accesses the contents of the vault using:

```
fun vaultQueryBy(  
    criteria: QueryCriteria,  
    paging: PageSpecification,  
    sorting: Sort): Vault.Page<T>
```

- This returns a page of the states matching the vault query criteria



# Attachments

- **CordaRPCOps** defines three methods for interacting with attachments:

`fun attachmentExists(id: SecureHash): Boolean`

`fun openAttachment(id: SecureHash): InputStream`

`fun uploadAttachment(jar: InputStream): SecureHash`





# Implementation

# Implementation

- One implementation of a node is **AbstractNode** (see [net.corda.node.internal.AbstractNode](#))
- **AbstractNode**'s **ServiceHub** is initialized in its **start()** method
- The **AbstractNode** class is itself abstract, and subclasses must provide implementations for (at least):
  - **val** log: Logger
  - **val** networkMapAddress: SingleMessageRecipient?
  - **val** serverThread: AffinityExecutor
  - **fun** makeMessagingService(): MessagingServiceInternal
  - **fun** startMessagingService(rpcOps: RPCOps)




# Implementation – Recording transactions

`AbstractNode` records transactions using  
`ServiceHubInternal.recordTransactions(notifyVault: Boolean, txs: Iterable<SignedTransaction>)`

This will:

1. Records the transaction in local storage
2. Notify the vault



# Implementation

- Corda is agnostic about how the available RPC operations are implemented
- One implementation is [`net.corda.node.internal.CordaRPCOpsImpl`](#)
- In [`CordaRPCOpsImpl`](#), transactions and attachments are stored as blobs in key-value mappings – see:
  - [`DBTransactionStorage`](#)
  - [`NodeAttachmentService`](#)





# The node in summary

- Users interact with a node using a set of RPC operations
- Internally, a node uses its ServiceHub to retrieve and record data, and to start flows
- **AbstractNode** is one implementation of the node architecture, but other implementations are possible



r3

Practical

# Designing the Network

- We are now going to move beyond the test DSLs and make this real:
  - We'll define the nodes on our network
  - We'll deploy (i.e. build) the nodes we've defined
  - We'll run the nodes

**1. CorDapp Design**

**2. State**

**3. Contract**

**4. Flow**

**5. Network**

- Defining nodes
- Deploying nodes
- Running nodes
- ✓ Checkpoint

**6. API**




# **Step 1 – Defining the Nodes**



# Network Layout

- We will use the following network design:



1. CorDapp Design

2. State

3. Contract

4. Flow

5. Network

- Defining nodes
- Deploying nodes
- Running nodes
- ✓ Checkpoint

6. API

# The DeployNodes Task

- Nodes are built using the Gradle build system
- Gradle builds are organized into tasks
- Nodes are created using the **deployNodes** task:

```
node {  
    name "O=Network Map Service,L=London,C=GB"  
advertisedServices = ["corda.notary.validating"]  
    p2pPort 10002  
    rpcPort 10003  
    cordapps = []  
}
```

- **deployNodes** extends the Cordform plugin:
  - <http://jcenter.bintray.com/net/corda/plugins/cordformation/>

1. CorDapp Design

2. State

3. Contract

4. Flow


5. Network

- Defining nodes
- Deploying nodes
- Running nodes
- ✓ Checkpoint

6. API



# **Step 2 – Deploying the Nodes**



# Deploying the Nodes

- The **deployNodes** task for the IOU CorDapp is defined in build.gradle
- The build.gradle file shows that our network is made up of 4 nodes:
  - The Network Map Service, who runs the notary and the network map
  - Three regular nodes:
    - ParticipantA
    - ParticipantB
    - ParticipantC
- Running the **deployNodes** task will use these configs to build actual nodes

1. CorDapp Design

2. State

3. Contract

4. Flow

5. Network

- Defining nodes
- **Deploying nodes**
- Running nodes
- ✓ Checkpoint

6. API



# Deploying the Nodes - Instructions



<b>Goal</b>	Deploy the nodes
<b>Where?</b>	The command line
<b>Steps</b>	<ol style="list-style-type: none"><li>1. From the command line, move to the root of your CorDapp directory</li><li>2. Deploy the nodes:<ul style="list-style-type: none"><li>• Unix: <a href="#">./gradlew deployNodes</a></li><li>• Windows: <a href="#">gradlew deployNodes</a></li></ul></li><li>3. Navigate to the build/nodes folder to view the deployed nodes</li></ol>
<b>Key Docs</b>	N/A

1. CorDapp Design

2. State

3. Contract

4. Flow

5. Network

- Defining nodes
- **Deploying nodes**
- Running nodes
- ✓ Checkpoint

6. API

r3.

# **Step 3 – Running the Nodes**



# The runnodes Script

- The deployed nodes are found in the build/nodes folder
- Each node folder contains a corda.jar file of our IOU CorDapp
- We can use the **runnodes** script to start up all the nodes

1. CorDapp Design

2. State

3. Contract

4. Flow

5. Network

- Defining nodes
- Deploying nodes
- **Running nodes**
- ✓ Checkpoint

6. API



# The runnodes Script - Instructions



<b>Goal</b>	Run the nodes
<b>Where?</b>	The command line
<b>Steps</b>	<ol style="list-style-type: none"><li>1. On the command line, navigate to the build/nodes folder to view the deployed nodes</li><li>2. Execute the <a href="#">runnodes</a> script, using:<ul style="list-style-type: none"><li>• Unix: <code>sh runnodes</code></li><li>• Windows: <code>runnodes</code></li></ul></li><li>3. All the nodes will start up</li><li>4. There should be seven terminal windows in all</li></ol>
<b>Key Docs</b>	N/A

r3.

# The Node Terminal Window

- The first four terminal windows show information on each node:



The terminal window displays a developer snapshot for Node A. It includes a small icon of a computer monitor with a grid pattern, a message from the node, and a log output.

"It's OK computer, I go to sleep after twenty minutes of inactivity too!"

--- DEVELOPER SNAPSHOT -----

```
Logs can be found in : /Users/joeldudley/Desktop/cordapp-tutorial/java-source/build/nodes/NodeA/logs
Database connection url is : jdbc:h2:tcp://10.21.0.137:50074/node
Node listening on address : localhost:10004
Loaded plugins : com.example.plugin.ExamplePlugin
Node started up and registered in 22.49 sec
```

1. CorDapp Design

2. State

3. Contract

4. Flow

5. Network


- Defining nodes
- Deploying nodes
- Running nodes
- ✓ Checkpoint

6. API



# The Webserver Terminal Window

- The remaining four windows show each node's webserver:



A terminal window displaying the logs for Node A. The window features a small network diagram icon in the top-left corner. The text output includes a purple quote: "It's not who you know, it's who you know knows what you know you know." Below this, a dashed line separates the quote from the log output. The log output itself is white text on a black background, starting with "--- DEVELOPER SNAPSHOT ---" and listing the path to the logs and the start time of the webserver.

```
It's not who you know, it's who you know  
knows what you know you know.  
--- DEVELOPER SNAPSHOT ---  
Logs can be found in : /Users/joeldudley/Desktop/cordapp-tutorial/java-source/build/nodes/NodeA/logs/web  
Starting as webserver: localhost:10005  
Webserver started up in 25.38 sec
```

1. CorDapp Design

2. State

3. Contract

4. Flow

5. Network


- Defining nodes
- Deploying nodes
- Running nodes
- ✓ Checkpoint

6. API



# **Checkpoint – Progress So Far**





# Our progress so far

- Our nodes are up and running
- But we still don't have a way to interact with them
- For that, we'll be building a simple REST API

r3.